

Presto

- An open source distributed SQL query engine for running interactive analytic queries against data sources of all sizes ranging from gigabytes to petabytes.
- Allows querying data where it lives, including Hive, Cassandra, relational databases or even proprietary data stores.
- A single Presto query can combine data from multiple sources, allowing for analytics across your entire organization.
- Originally developed and outsourced by Facebook
- Airbnb developed and outsourced Airpal web interface for Presto

Presto Query

```
SELECT a, b, c  
FROM ( SELECT a, MAX(b) AS b, MIN(c) AS c  
      FROM tbl GROUP BY a ) tbl_alias
```

Presto Query

```
WITH tbl_alias AS (SELECT a, MAX(b) AS b, MIN(c) AS c  
                    FROM tbl GROUP BY a)  
SELECT a, b, c FROM tbl_alias
```

Presto Query

```
WITH tbl1 AS (SELECT a, MAX(b) AS b, MIN(c) AS c  
               FROM tbl GROUP BY a),  
tbl2 AS (SELECT a, AVG(d) AS d  
         FROM another_tbl GROUP BY a)  
SELECT tbl1.*, tbl2.* FROM tbl1 JOIN tbl2 ON tbl1.a = tbl2.a
```